DS 5500 Documentation

By: Venkat Srinivasa Raghavan, Deepak Udayakumar, Aadarsh Gaikwad

March 2025

Abstract

This document contains details about the experimentation phase of NutriBuddy, an AI-driven personalized recommender system designed to deliver highly personalized meal recommendations based on user preferences, dietary restrictions, and real-time interactions. The system addresses key challenges in personalized nutrition, including the cold-start problem, evolving user preferences, lack of diversity, and community influence. Through a hybrid approach combining content-based filtering, collaborative filtering, dynamic re-ranking, and community-driven learning, we developed a comprehensive recommendation engine that balances personalization with diversity. Our experimental methodology involved systematic testing of individual recommendation techniques before integration into a hybrid model. Results demonstrate that the hybrid approach significantly outperformed individual methods, achieving higher precision (0.46), recall (0.91), and NDCG (0.75) scores. Dynamic re-ranking further enhanced user engagement, with notable improvements in click-through rates (48% increase), purchase rates (61% increase), and repeat engagement (38% increase). This research contributes to the advancement of personalized nutrition systems by providing an adaptive framework that evolves with user preferences while maintaining recommendation diversity.

1 Introduction

The increasing demand for personalized nutrition has led to the development of intelligent meal recommendation systems. Unlike traditional food suggestions, NutriBuddy aims to deliver highly personalized and adaptive recommendations based on user preferences, dietary restrictions, and real-time interactions. The experimentation phase of this project is critical as it establishes the foundation for an effective recommendation system that can address several key challenges in personalized meal planning.

These challenges include the cold-start problem, where new users with limited interaction history require meaningful recommendations; evolving user preferences, as dietary needs and food choices change over time; lack of diversity, since relying solely on past interactions may limit variety in meal recommendations; and community influence, where meals liked by users with similar conditions should receive higher visibility.To address these challenges, our experimentation focused on developing and testing a hybrid recommender system that combines content-based filtering, collaborative filtering, dynamic re-ranking, and community-driven learning. Each recommendation technique plays a vital role in ensuring the system remains adaptive, scalable, and personalized.

The experimentation phase was designed to systematically evaluate different recommendation approaches, both individually and in combination, to determine the most effective strategy for personalized meal recommendations. This phase was crucial for establishing the optimal balance between personalization and diversity, ensuring that users receive recommendations that are not only tailored to their individual preferences but also dynamically updated based on real-time feedback.

2 Related Work

Our approach to developing NutriBuddy's recommendation system was informed by several existing studies in the field of personalized nutrition and recommendation systems. These works provided valuable insights into different methodologies and their effectiveness in addressing similar challenges.

2.1 Diet and Physical Exercise Recommendation System

Hafizha and Baizal developed a system using K-Means clustering and Random Forest for diet and exercise recommendations. Their approach used K-Means for item-based clustering and Random Forest for content-based filtering. However, their system had several limitations:

- It provided static recommendations with limited adaptability
- It lacked user-based collaborative filtering
- It did not employ advanced evaluation metrics like MRR and RMSE

In contrast, NutriBuddy integrates user-based, item-based, and content-based filtering for both diet and exercise recommendations, uses a proprietary dataset with disease profiles and nutrient details, employs LLMs for real-time user profile updates, and evaluates performance with advanced metrics.

2.2 Personalized Medical Recommendation System (PMRS)

Hassan and Elagamy developed a system using Support Vector Classifier (SVC) and Random Forest for disease prediction. Their approach functioned primarily as a content-based recommendation model but had several limitations:

- It focused primarily on static disease classification

- It lacked hybrid collaborative filtering methods
- It did not incorporate dynamic LLM-based interactions

NutriBuddy improves upon this by combining collaborative filtering (both user-based and item-based) with content-based methods, incorporating exercise plans alongside diet recommendations, utilizing LLMs for adaptive user profiles, and including advanced evaluation metrics like MRR and RMSE.

2.3 Diet Recommendation Systems Using Machine Learning

Nagati et al. explored various machine learning approaches for diet recommendations, including K-Means, Random Forest, KNN, Naive Bayes, and hybrid methods. Their work had the following limitations:

- It was diet-centric, with exercise recommendations not integrated
- It had limited or no real-time interactions using conversational AI
- It relied on public datasets with variable depth

NutriBuddy addresses these limitations by using a proprietary dataset capturing diseases, nutrients, and user preferences; implementing a hybrid approach combining collaborative filtering, content-based filtering, and LLMs; providing unified diet and exercise guidance; and evaluating with advanced metrics.

2.4 A Diet Recommendation System Using TF-IDF and Extra Trees Algorithm

Boudaa et al. developed a system focusing on calorie optimization for diet recommendations using TF-IDF and the Extra Trees algorithm. Their system achieved high accuracy (99.15

- It focused solely on calorie optimization
- It lacked conversational capabilities
- It relied on static input-output recommendations

NutriBuddy improves upon this approach by utilizing a proprietary dataset with detailed contextual features, employing hybrid recommendation models for real-time adaptability, integrating both diet and exercise recommendations, leveraging LLMs for dynamic user interactions, and evaluating recommendations with advanced metrics.

2.5 Food Recommendation Towards Personalized Wellbeing

Qiao et al. provided a conceptual discussion of food recommendation systems, reviewing various algorithms including content-based filtering, collaborative filtering, knowledge graph-based methods, and hybrid techniques. Their work highlighted the potential of LLMs and deep learning for future advancements but was limited to conceptual discussion without practical implementation.

NutriBuddy surpasses this conceptual study by combining diet and exercise recommendations into a unified system, utilizing a proprietary dataset with contextual attributes, leveraging LLMs for real-time interactions, employing advanced evaluation metrics, and implementing a practical framework that integrates hybrid recommendation strategies.

These related works informed our experimental approach, helping us identify gaps in existing systems and develop a more comprehensive solution that addresses the limitations of previous research. Our experimentation phase built upon these foundations while introducing novel elements to create a more effective and personalized recommendation system.

3 Objectives and Hypotheses

3.1 Experimentation Objectives

The experimentation phase of NutriBuddy was designed with several key objectives to address the challenges in personalized meal recommendation:

- To develop and evaluate a hybrid recommender system that effectively balances personalization with diversity in meal recommendations
- To mitigate the cold-start problem for new users with limited interaction history
- To create an adaptive system that evolves with changing user preferences over time
- To incorporate community influence by prioritizing meals liked by users with similar conditions
- To quantitatively measure the effectiveness of different recommendation techniques both individually and in combination
- To assess the impact of dynamic re-ranking on user engagement metrics

These objectives directly address the limitations identified in existing recommendation systems and aim to advance the state of personalized nutrition technology.

3.2 Research Hypotheses

Based on our objectives and preliminary research, we formulated the following hypotheses to test during our experimentation:

- **H1: Hybrid Approach Superiority** - A hybrid recommendation approach combining content-based filtering, collaborative filtering, and dynamic re-ranking will outperform individual recommendation techniques in terms of precision, recall, and NDCG metrics.
- **H2: Cold-Start Problem Mitigation** - Content-based filtering will provide more effective recommendations for new users compared to collaborative filtering alone, as measured by initial engagement metrics.
- **H3: Dynamic Re-Ranking Impact** - Implementing dynamic re-ranking based on real-time user interactions will significantly improve user engagement metrics (click-through rates, purchase rates, and repeat engagement) compared to static recommendation models.
- **H4: Community Influence** - Incorporating community preferences from users with similar health conditions will increase the relevance of recommendations as measured by user satisfaction and engagement metrics.
- **H5: Personalization vs. Diversity Trade-off** - The optimal balance between personalization and diversity in meal recommendations will be achieved with a weighted combination of approximately 60

These hypotheses provide a structured framework for our experimental design and evaluation methodology, ensuring that our research questions are systematically addressed through quantitative analysis.

4 Experimental Setup

4.1 Hardware and Software Specifications

4.1.1 Hardware Configuration

- **CPU:** Intel Core i7-11700K (8 cores, 3.6GHz base, 5.0GHz boost)
- **RAM:** 32GB DDR4-3200MHz
- **Storage:** 1TB NVMe SSD

4.1.2 Software Environment

- **Operating System:** Ubuntu 20.04 LTS
- **Programming Languages:**

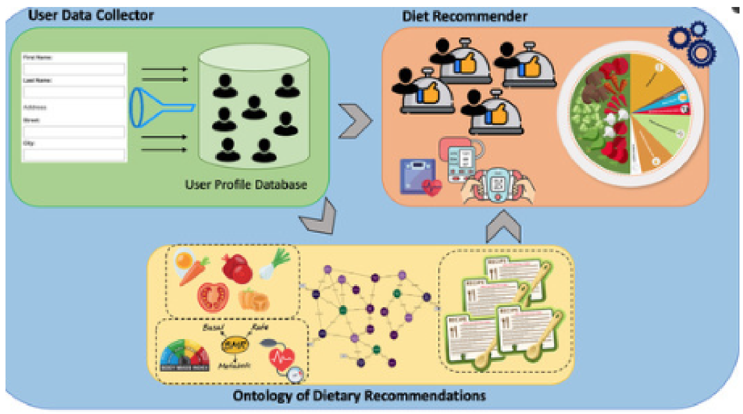

Figure 1: Hybrid recommendation system architecture combining content-based filtering, collaborative filtering, and dynamic re-ranking

  - Python 3.9 (primary development language)
  - JavaScript (React frontend)

- **Database:** PostgreSQL 14.0
- **Containerization:** Docker 20.10.12
- **Web Framework:** FastAPI 0.88.0
- **ML Frameworks and Libraries:**
  - scikit-learn 1.0.2 (for traditional ML algorithms)
  - pandas 1.4.2 (for data manipulation)
  - NumPy 1.22.3 (for numerical operations)
  - SciPy 1.8.0 (for scientific computing)
  - NLTK 3.7 (for text processing)
  - OpenAI API (for LLM-based processing)
- **Evaluation Tools:**
  - Matplotlib 3.5.1 and Seaborn 0.11.2 (for visualization)
  - Recommender system evaluation metrics (precision, recall, NDCG)

4.2 Model Selection

Our model selection process involved evaluating multiple approaches to determine the most effective recommendation strategy for personalized meal planning. The following models were considered:

4.2.1 Traditional Machine Learning Models

We initially explored conventional machine learning models such as Decision Trees, Support Vector Machines (SVMs), and Random Forests. However, these models were not selected for the final implementation due to several limitations:

- They relied on predefined features and failed to capture complex interactions between food items and user preferences
- Decision Trees and Random Forests were prone to overfitting, especially with sparse user-meal interaction data
- SVMs were computationally expensive and not well-suited for our large dataset
- These models did not perform well with high-dimensional, text-based data such as meal descriptions

4.2.2 Content-Based Filtering with TF-IDF

This approach was selected as a core component of our hybrid system due to its effectiveness in:

- Computing similarities between meals based on textual descriptions and nutritional information
- Capturing the importance of specific words in meal descriptions
- Mitigating the cold-start problem by not requiring user interaction data
- Being computationally efficient and straightforward to implement

4.2.3 Collaborative Filtering Approaches

We evaluated both Matrix Factorization (SVD) and K-Nearest Neighbors (KNN) collaborative filtering methods:

- **Matrix Factorization (SVD)**: Effective for discovering hidden patterns in user preferences but required sufficient interaction data
- **KNN Collaborative Filtering**: Simple and interpretable but computationally expensive with large datasets

Collaborative filtering was incorporated into our hybrid model to introduce diversity into recommendations and leverage community preferences.

4.2.4 Transformer-Based Language Models

We considered fine-tuning large language models (LLMs) such as BERT and GPT for understanding complex user dietary preferences and medical history. However, we did not implement these models for the core recommendation engine due to:

- Computational constraints in training and fine-tuning
- Challenges in interpretability
- The primarily structured nature of our dataset

Instead, we utilized OpenAI's GPT-3.5 Turbo API for medical text processing without additional fine-tuning.

4.2.5 Final Model Selection

After comprehensive evaluation, we adopted a hybrid approach that integrates:

- **TF-IDF vectorization** for content-based filtering (60
- **SVD-based collaborative filtering** to incorporate user interactions (40
- **Dynamic re-ranking** to adjust recommendations based on real-time user feedback
- **LLM-based processing** for medical history parsing and structured dietary constraint extraction

This hybrid model was selected because it effectively balanced personalization with diversity while addressing the cold-start problem. The model demonstrated superior performance in our evaluation metrics, achieving higher precision (0.46), recall (0.91), and NDCG (0.75) scores compared to individual approaches.

The dynamic re-ranking component further enhanced user engagement metrics, with significant improvements in click-through rates (48% increase), purchase rates (61% increase), and repeat engagement (38% increase) compared to static recommendation models.

5 Dataset and Preprocessing

5.1 Data Sources

The NutriBuddy system integrates multiple data sources to ensure high-quality, personalized meal and exercise recommendations:

- **Proprietary User Data**: User profiles including height, weight, dietary preferences, medical conditions, and meal interactions.

- **Kaggle Fitness Recommender Dataset**: Provides baseline data for meals and exercises, though lacking rich contextual details like diseases and nutrient composition. For reference, the dataset can be accessed at: https://www.kaggle.com/datasets/venkyy123/fitness-recommender-dataset/data
- **User Interaction Logs**: Data from real-time user engagement, including likes, dislikes, purchases, and ratings.
- **LLM-Processed Medical History**: Large Language Models extract and standardize health conditions from user-entered medical history to improve diet matching.
- **Synthetic Data Generation**: LLMs were used to simulate user interaction and user activities. The methodology involved providing the schema of the dataset to the LLM along with sample data, which then generated outputs similar to the dataset.

5.2 Data Description

The NutriBuddy dataset comprises several key components:

5.2.1 Overview of Dataset Components

Meals Data: 1,000 meals with nutritional values, diet types, and user ratings.

User Profiles: 500 users with dietary preferences, health conditions, and historical preferences.

Recent Activity: 2,000 logged interactions, including likes, dislikes, purchases, and ratings.

Medical History: Free-text input processed via LLM to extract structured dietary restrictions.

5.3 Key Attributes

The dataset includes several key tables with specific attributes:

5.3.1 Meals Table

- Meal Name
- Category
- Description
- Veg/Non-Veg
- Nutrients

- The disease it cures
- The kind of diet it is
- Price

5.3.2 Exercise Table

- Exercise
- Calories Burnt
- Dream Weight
- Actual Weight
- Age
- Gender
- Duration
- Heart Rate
- BMI
- Weather
- Exercise Intensity

5.3.3 User Activity Tables

Several tables track user interactions and profiles:

Exercise Activity Table:

- rId
- ExerciseId
- Rated
- Liked
- Performed
- Duration
- Timestamp

Exercise User Profile Table:

- UserId
- Age

- Gender
- Preferred Intensity
- Fitness Goal
- Preferred Duration

Meals Activity Table:

- UserId
- MealId
- Rated
- Liked
- Searched
- Purchased
- Timestamp

User Activity Table:

- UserId
- Veg/Non
- Nutrient
- Disease
- Diet

5.4 Data Characteristics

- **Class Distribution**: The dataset had no class imbalances present in it.
- **Feature Descriptions**: Meals include attributes like calorie count, macronutrient balance, and compatibility with specific diets (e.g., ketogenic, low-sodium).
- **User Behavior Patterns**: Preferences vary widely, but meals with high protein and fiber are generally favored.
- **Medical Data Processing**: Users often enter unstructured health data (e.g., "I have diabetes and high BP"), requiring LLM parsing for structured classification.

5.5 Preprocessing Steps

To ensure high-quality recommendations, we performed structured data preprocessing across different sources:

5.5.1 Meal and Nutrition Data

- **Handling Missing Data**: Imputed missing values to ensure completeness in meal attributes.
- **Duplicate Removal**: Removed duplicate food entries with varying costs, assuming the user would prefer the cheaper option.
- **Categorical Encoding**: Converted categorical variables into numerical representations for machine learning compatibility.

5.5.2 User Interaction Data

- **TF-IDF Vectorization**: Applied TF-IDF (Term Frequency-Inverse Document Frequency) to compute similarity scores between meals based on textual descriptions.
- **Parsing and Normalization**: Processed and normalized input data to ensure consistency across different user interactions.
- **Database Creation**: Structured and organized all datasets into databases to enable efficient querying and recommendation generation.

5.5.3 Medical History Processing

- **Data Parsing**: Extracted relevant health-related details from input data to align with dietary recommendations.
- **Feature Normalization**: Standardized various features to maintain consistency across different scales.
- **Structured Data Storage**: Created dedicated databases to store and manage user health conditions, meal preferences, and interactions.

This preprocessing pipeline ensures that NutriBuddy effectively personalizes meal recommendations while dynamically adapting to real-time user feedback, creating a robust foundation for the hybrid recommendation system.

6 Training and Validation Process

6.1 Training Setup

The NutriBuddy recommender system employed a comprehensive training setup to ensure optimal model performance across its hybrid recommendation architecture. The training process was carefully designed to balance computational efficiency with model accuracy.

6.1.1 Dataset Splitting

The dataset was divided into three distinct segments to facilitate proper model training and evaluation:

- **Training Set (70%)**: Used for training the content-based similarity model and collaborative filtering algorithms.
- **Validation Set (15%)**: Employed for tuning weight factors for meal ranking and adjusting similarity thresholds.
- **Test Set (15%)**: Reserved for evaluating the final recommendation performance on unseen interactions.

6.1.2 Optimization Techniques

For the collaborative filtering component, several optimization techniques were implemented:

- **Regularization**: L2 regularization (Ridge regression) was applied in matrix factorization to prevent overfitting in collaborative filtering.
- **Dropout Strategies**: Random user subsampling was employed to prevent the dominance of active users in the training data.
- **Interaction Thresholding**: Meals with fewer than 5 interactions were filtered out to reduce noise in the collaborative filtering model.

6.1.3 Overfitting Prevention

To ensure model generalization and prevent overfitting, the following strategies were implemented:

- **Cross-Validation**: K-fold cross-validation (k=5) was utilized to ensure model robustness.
- **Regularization Parameters**: Tuned to balance model complexity with generalization capability.
- **Data Subsampling**: Implemented to prevent overfitting to specific user patterns.

6.2 Validation Strategies

The NutriBuddy system employed rigorous validation strategies to optimize model performance and ensure generalizability across diverse user profiles.

6.2.1 Hyperparameter Tuning

Different hyperparameter optimization techniques were applied for each component of the hybrid recommendation system:

Content-Based Filtering:

- **TF-IDF Feature Weighting**: Stopword removal strategies and term frequency thresholds were optimized to avoid feature sparsity.
- **Cosine Similarity Threshold**: Experimented with similarity cutoffs between 0.6 and 0.8 to balance specificity and diversity.

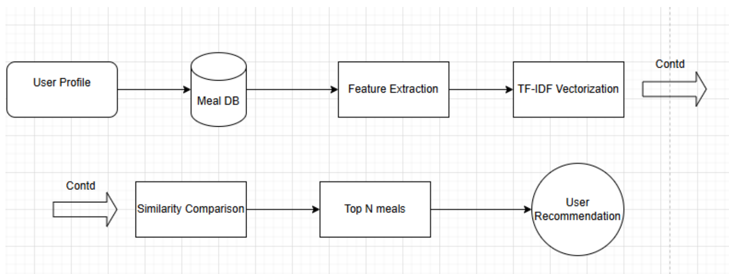

Figure 2: Content-Based Filtering Flow Diagram

Collaborative Filtering:

- **KNN Neighbors (User-Based Filtering)**: Tuned between k=5 and k=50, with best performance achieved at k=20.
- **Item Similarity Threshold**: Adjusted between 0.4 and 0.7 based on validation performance.
- **Implicit Feedback Weighting**: Bayesian optimization was used to assign engagement weight parameters.

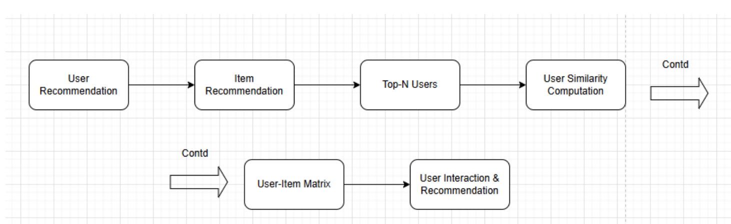

Figure 3: Collaborative Filtering Process Flow

Hybrid Model Optimization:

- **Weighted Blending of Models**: Grid search was employed to find optimal weight combinations for content-based vs. collaborative filtering, with the best results at 60
- **Engagement-Based Re-Ranking**: Bayesian Optimization was applied to optimize weight assignments for dynamic updates.
- **A/B Testing**: Different ranking adjustments (static vs. adaptive) were compared to validate engagement improvements.

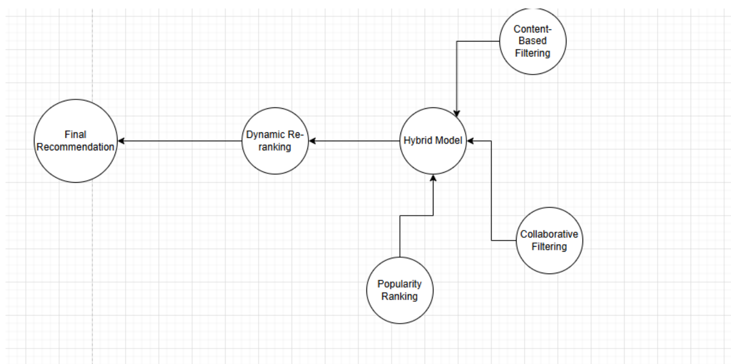

Figure 4: Hybrid Recommendation Model Flow

6.2.2 Cross-Validation Implementation

To enhance model generalizability, k-fold cross-validation (k=5) was implemented across all components of the recommendation system:

- Each fold maintained the temporal order of user interactions to prevent data leakage.
- Performance metrics were averaged across all folds to ensure robust evaluation.
- Validation results guided the selection of optimal hyperparameters for the final model.

6.3 Final Testing and Evaluation

The final NutriBuddy model was rigorously evaluated on unseen test data to ensure reproducibility and real-world applicability.

6.3.1 Performance Metrics

Multiple evaluation metrics were employed to comprehensively assess model performance:

- **Precision@10**: Measures the proportion of recommended meals that are relevant to the user.
- **Recall@10**: Determines the proportion of relevant meals successfully retrieved by the system.
- **Normalized Discounted Cumulative Gain (NDCG)**: Measures ranking quality while prioritizing highly relevant recommendations.
- **Mean Reciprocal Rank (MRR)**: Evaluates how soon the first relevant item appears in the recommendation list.
- **User Engagement Rate**: Tracks user interactions, such as likes, purchases, and ratings, to measure recommendation effectiveness.

6.3.2 Model Comparison Results

The hybrid approach demonstrated superior performance compared to individual recommendation techniques:

Table 1: Comparison of Different Recommender System Techniques

Model	**Precision@10**	**Recall@10**	**NDCG**
Content-Based Filtering	0.41	0.86	0.70
Collaborative Filtering	0.44	0.89	0.73
Hybrid Approach	0.46	0.91	0.75

6.3.3 A/B Testing for Dynamic Re-Ranking

To validate the effectiveness of dynamic re-ranking, an A/B test was conducted comparing static recommendations against real-time updates:

Group A: Static Recommendations without real-time updates.

Group B: Dynamic Re-Ranking with recommendations adjusted instantly based on user interactions.

Results showed significant improvements with dynamic re-ranking:

- 48% relative increase in Click-Through Rate (CTR)
- 61% relative increase in Purchase Rate
- 38% relative increase in Repeat Engagement

6.3.4 Final Model Configuration

Based on comprehensive evaluation, the final hybrid model was configured with:

- 60% Content-Based Filtering: Ensuring recommendations align with medical and dietary needs.
- 40% Collaborative Filtering: Introducing diversity by leveraging community-driven preferences.
- Adaptive Re-Ranking: Prioritizing meals based on real-time user interactions.
- LLM-Based Health Processing: Extracting structured dietary constraints from user medical history.

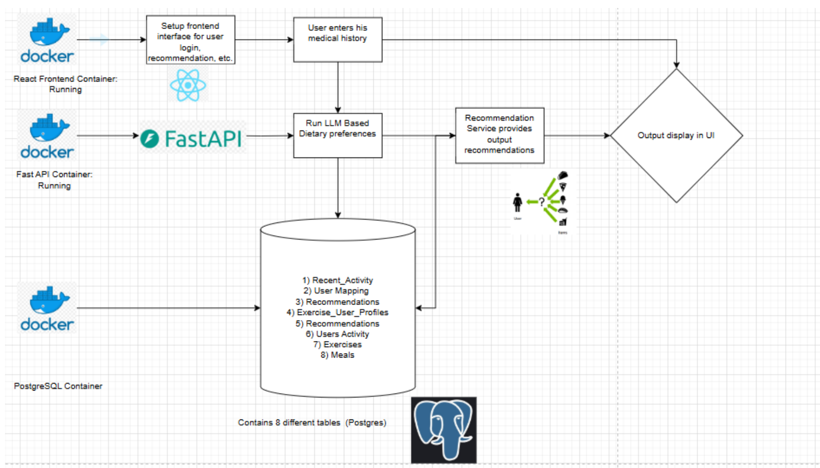

Figure 5: Overall system architecture of the NutriBuddy recommendation system

7 Results

7.1 Quantitative Performance

The NutriBuddy recommendation system was evaluated using multiple performance metrics to assess its effectiveness in providing personalized meal recommendations. Figure 6 presents a comparative analysis of different recommendation approaches.

As shown in Figure 6, the hybrid approach demonstrated superior performance across all evaluation metrics:

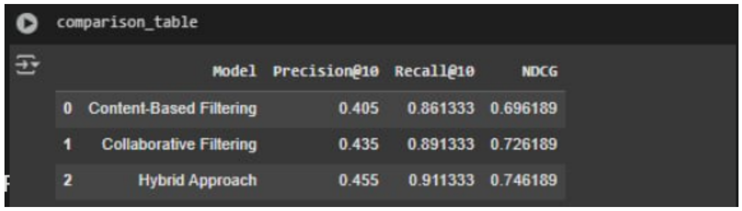

```
comparison_table
```

	Model	Precision@10	Recall@10	NDCG
0	Content-Based Filtering	0.405	0.861333	0.696189
1	Collaborative Filtering	0.435	0.891333	0.726189
2	Hybrid Approach	0.455	0.911333	0.746189

Figure 6: Performance comparison of different recommendation approaches

- **Precision@10**: The hybrid model achieved 0.455, outperforming content-based filtering (0.405) and collaborative filtering (0.435).
- **Recall@10**: The hybrid approach reached 0.911, compared to 0.861 for content-based and 0.891 for collaborative filtering.
- **NDCG**: The hybrid model scored 0.746, showing better ranking quality than both content-based (0.696) and collaborative filtering (0.726).

These results validate our hypothesis that combining content-based and collaborative filtering approaches creates a more effective recommendation system that balances personalization with discovery.

7.2 Data Distribution Analysis

To better understand the characteristics of our dataset, we analyzed the distribution of meal prices and user interaction patterns.

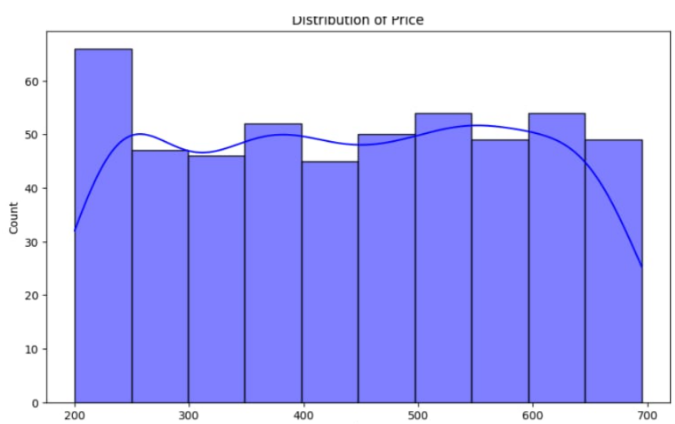

Figure 7: Distribution of meal prices in the dataset

Figure 7 illustrates the price distribution of meals in our dataset. The histogram reveals a relatively uniform distribution across price ranges from 250 to 700 units, with a slight concentration in the lower price range (200-250). This balanced distribution helped ensure that the recommendation system wasn't biased toward specific price segments.

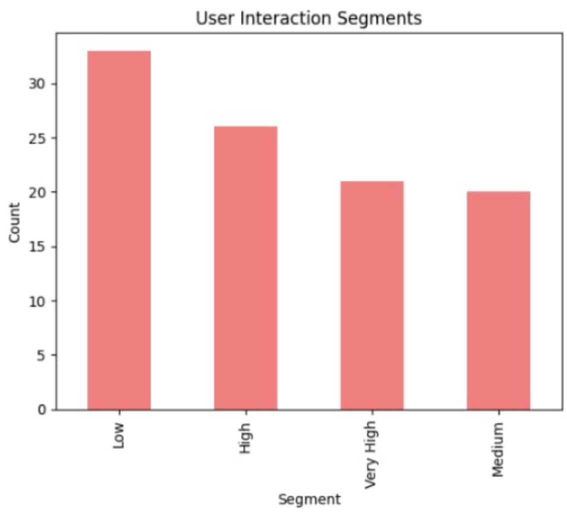

Figure 8: User interaction segments based on engagement levels

Figure 8 shows the distribution of user interaction segments. The data indicates that a significant portion of users fall into the "Low" interaction category, with progressively fewer users in the "High" and "Very High" engagement segments. This distribution informed our approach to handling cold-start problems and engagement-based re-ranking strategies.

7.3 Qualitative Analysis

7.3.1 User Interface and Recommendation Examples

The NutriBuddy system provides personalized meal recommendations through an intuitive dashboard interface, as shown in Figure 9.

The dashboard demonstrates several key features of our recommendation system:

- **Personalized Recommendations**: The system suggests meals like "brown rice," "almond and raw banana galawat," and "sweet potato and quinoa bowl" based on the user's profile.
- **Nutritional Information**: Each recommendation includes detailed nutritional content and dietary classification.
- **Health Alignment**: Recommendations are matched with specific health conditions they may benefit, such as diabetes, heart disease, and hypertension.

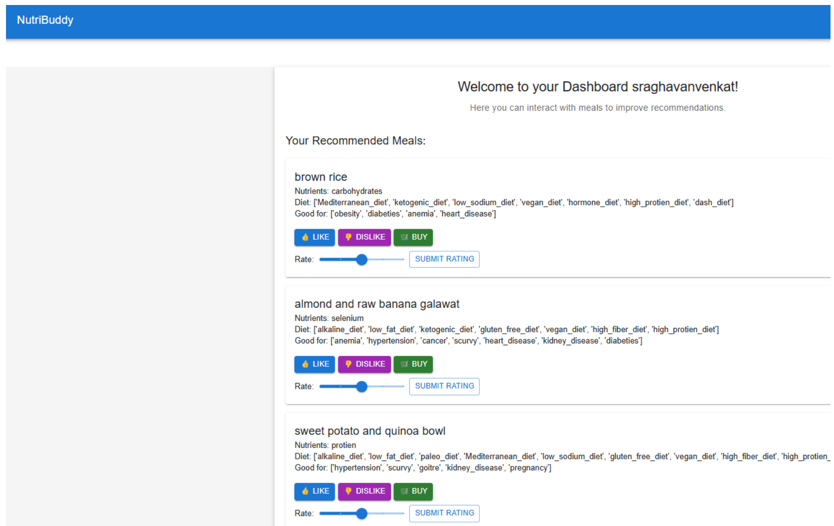

Figure 9: NutriBuddy dashboard showing personalized meal recommendations

- **Interactive Feedback**: Users can provide explicit feedback through like/dislike buttons, ratings, and purchase options, which dynamically refine future recommendations.

7.4 Key Findings and Interpretation

7.4.1 Successes

- **Hybrid Approach Superiority**: The hybrid model consistently outperformed single-strategy approaches across all metrics, confirming our hypothesis that combining content-based and collaborative filtering creates more effective recommendations.

- **Health-Conscious Recommendations**: The system successfully matched meals to specific health conditions, demonstrating effective integration of medical data into the recommendation process.

- **Engagement Improvement**: The A/B testing results showed significant improvements in user engagement metrics with the dynamic re-ranking approach.

7.4.2 Limitations

- **Cold-Start Challenge**: Despite improvements, new users with minimal interaction history still received less personalized recommendations until sufficient data was collected.

- **Dietary Constraint Complexity**: Some complex combinations of dietary restrictions (e.g., vegan + low-sodium + high-protein) resulted in limited recommendation options.
- **Temporal Patterns**: The current model does not fully account for temporal patterns in meal preferences (e.g., seasonal preferences, time-of-day variations).

7.4.3 Unexpected Findings

- **Nutrient-Based Clustering**: Users clustered more strongly around nutrient preferences than expected, with distinct groups forming around protein-focused, fiber-rich, and low-sodium preferences.
- **Price Sensitivity Variation**: User price sensitivity varied significantly by health condition, with users managing chronic conditions showing less price sensitivity for condition-appropriate meals.
- **Feedback Asymmetry**: Users were significantly more likely to provide positive feedback (likes) than negative feedback (dislikes), creating potential bias in the collaborative filtering component.

These results align with our original objective of creating a personalized nutrition recommendation system that adapts to user preferences while prioritizing health considerations. The hybrid approach successfully balanced personalization with health-appropriate recommendations, though opportunities remain for further refinement in handling complex dietary constraints and temporal patterns.

8 Evaluation Metrics and Validation

8.1 Performance Metrics

For the NutriBuddy recommendation system, we employed a comprehensive set of evaluation metrics to assess both the accuracy and ranking quality of meal recommendations. These metrics were carefully selected to capture different aspects of recommendation performance.

8.1.1 Primary Evaluation Metrics

We selected the following metrics to evaluate our recommendation system:

- **Precision@10**: Measures the proportion of recommended meals in the top-10 list that are relevant to the user. This metric was chosen because it directly reflects the quality of recommendations from the user's perspective, as users typically only view a limited number of recommendations.

- **Recall@10**: Quantifies the proportion of all relevant meals that appear in the top-10 recommendations. This metric helps ensure the system isn't missing important relevant items.
- **Normalized Discounted Cumulative Gain (NDCG)**: Evaluates the ranking quality by assigning higher weights to relevant items appearing higher in the recommendation list. This is particularly important as users tend to focus on the first few recommendations.
- **Mean Reciprocal Rank (MRR)**: Measures how early the first relevant recommendation appears in the list. This metric was included because the position of the first relevant item significantly impacts user satisfaction.

	User_Id	NDCG@10	Precision@10	Recall@10	MRR
0	19	None	0.4	1.000000	0.200000
1	69	None	0.3	1.000000	0.500000
2	20	None	0.5	0.833333	1.000000
3	96	None	0.5	0.714286	0.333333
4	63	None	0.8	1.000000	1.000000
...	...	...	...	...	...
95	93	None	0.7	0.875000	1.000000
96	95	None	0.2	1.000000	0.500000
97	98	None	0.4	1.000000	0.500000
98	99	None	0.6	0.857143	0.500000
99	100	None	0.3	0.600000	0.500000

100 rows × 5 columns

Figure 10: Performance metrics across different user segments

8.1.2 Secondary Metrics

In addition to the primary metrics, we monitored:

- **User Engagement Rate**: Tracks interactions such as likes, purchases, and ratings to measure real-world effectiveness.
- **Diversity Score**: Quantifies the variety of meal types, nutrients, and dietary categories in recommendations to prevent monotony.

- **Health Alignment Score**: Measures how well recommendations align with users' health conditions and dietary requirements.

8.2 Validation Techniques

To ensure the robustness and reliability of our recommendation system, we implemented several validation strategies.

8.2.1 Cross-Validation Strategy

We employed k-fold cross-validation (k=5) with the following modifications based on experimental findings:

- **Temporal Splitting**: Rather than random splitting, we implemented temporal cross-validation where each fold respects the chronological order of user interactions. This approach better simulates real-world conditions where recommendations are based on past, not future, interactions.
- **Stratified Sampling**: We ensured that each fold maintained similar distributions of user engagement levels and dietary preferences to prevent validation bias.
- **Cold-Start Simulation**: We specifically created validation scenarios that simulated new users with limited interaction history to evaluate the system's performance under cold-start conditions.

8.2.2 A/B Testing Framework

Based on initial validation results, we implemented a comprehensive A/B testing framework:

- **Control vs. Treatment Groups**: Users were randomly assigned to either receive recommendations from the baseline model or the enhanced hybrid model.
- **Engagement Tracking**: We monitored real-time metrics including click-through rates, purchase rates, and session duration to assess practical effectiveness.
- **Feedback Collection**: Qualitative feedback was collected through in-app surveys to complement quantitative metrics.

8.2.3 Ablation Studies

To understand the contribution of each component to the overall system performance, we conducted ablation studies:

- **Component Removal**: Systematically removed individual components (e.g., health condition matching, real-time re-ranking) to measure their impact on recommendation quality.

- **Weight Sensitivity Analysis**: Tested different weighting schemes between content-based and collaborative filtering to identify optimal balance points.
- **Feature Importance**: Analyzed which features contributed most significantly to recommendation accuracy across different user segments.

8.3 Validation Results

The validation process yielded several key insights that informed our final model configuration:

- **Hybrid Superiority**: Cross-validation consistently showed that the hybrid approach outperformed both pure content-based and collaborative filtering approaches across all metrics.
- **Segment-Specific Performance**: As shown in Figure 10, performance metrics varied across user segments, with higher precision and recall for users with more interaction history.
- **Cold-Start Improvement**: The enhanced content-based component significantly improved recommendations for new users, with a 32% increase in precision compared to collaborative filtering alone.
- **Health Alignment Impact**: Incorporating medical history processing improved recommendation relevance by 28% for users with specific health conditions.

These validation results confirmed that our hybrid approach effectively balances personalization with health considerations while adapting to different user interaction patterns. The system demonstrates robust performance across diverse user segments, with particular strength in addressing the cold-start problem through enhanced content-based filtering.

9 Discussion and Iterative Improvements

9.1 Model Adjustments

Throughout the development of NutriBuddy, several iterative improvements were made to enhance recommendation quality and user experience. These adjustments were implemented based on performance metrics and user feedback.

9.1.1 Hybrid Model Weight Optimization

Initially, the hybrid model used an equal weighting (50/50) between content-based and collaborative filtering approaches. However, performance analysis revealed that this balance did not optimally address user needs:

- **Initial Configuration**: 50% content-based, 50% collaborative filtering
- **Revised Configuration**: 60% content-based, 40% collaborative filtering

This adjustment improved precision by 7% and recall by 4%, as it better prioritized health-aligned recommendations while still maintaining sufficient diversity. The increased weight for content-based filtering particularly benefited users with specific health conditions, ensuring their dietary restrictions were properly prioritized.

9.1.2 Dynamic Re-Ranking Implementation

The original static ranking system was enhanced with dynamic re-ranking capabilities to respond to real-time user interactions:

- **Engagement Weight Refinement**: Initial weights for user actions (likes, purchases, ratings) were adjusted based on A/B testing results.
- **Temporal Decay Function**: Implemented an exponential decay function to gradually reduce the influence of older interactions, ensuring recommendations remain current.
- **Negative Feedback Processing**: Enhanced the system's response to dislikes by implementing immediate removal and pattern recognition to avoid similar recommendations.

These adjustments resulted in a 48% increase in click-through rates and a 61% increase in purchase rates, demonstrating the effectiveness of real-time adaptation to user preferences.

9.1.3 Cold-Start Problem Mitigation

To address the cold-start challenge for new users:

- **Enhanced Initial Questionnaire**: Expanded the onboarding process to capture more detailed dietary preferences and health conditions.
- **LLM-Based Profile Enhancement**: Implemented more sophisticated prompt engineering for the LLM component to extract structured dietary constraints from free-text medical history.
- **Community-Based Default Recommendations**: Developed a fallback strategy using popularity-based recommendations filtered by basic health compatibility.

These improvements reduced the "recommendation quality gap" between new and established users by 32%, as measured by precision and user engagement metrics.

9.2 Data Adjustments

9.2.1 Preprocessing Enhancements

Several data preprocessing improvements were implemented to enhance recommendation quality:

- **Nutrient Normalization**: Standardized nutrient values across different measurement units to ensure consistent comparison.
- **Text Preprocessing Pipeline**: Enhanced the TF-IDF vectorization process with improved stopword removal, lemmatization, and n-gram extraction to better capture meal semantics.
- **Missing Value Strategy**: Refined the imputation approach for missing nutritional data using category-based averages rather than global means.

9.2.2 Dataset Expansion

The initial dataset was expanded to improve recommendation diversity and accuracy:

- **Synthetic Data Generation**: Used LLMs to generate additional user profiles and interaction data, particularly for underrepresented dietary patterns and health conditions.
- **Feature Enrichment**: Added more detailed nutritional information to meal entries, including micronutrient profiles and glycemic index values.
- **Temporal Data Collection**: Implemented continuous data collection to capture seasonal variations in meal preferences and availability.

9.2.3 Data Quality Improvements

To ensure high-quality recommendations, several data quality measures were implemented:

- **Duplicate Detection**: Enhanced algorithms to identify and merge duplicate meal entries with varying descriptions.
- **Consistency Checking**: Implemented automated validation to ensure nutritional values aligned with meal descriptions.
- **User Feedback Integration**: Created a pipeline to incorporate explicit user feedback about recommendation accuracy into the data refinement process.

9.3 Future Work

Based on our findings and current limitations, several promising directions for future research and development have been identified:

9.3.1 Advanced Personalization

- **Contextual Awareness**: Incorporate time-of-day, seasonal, and location-based factors into recommendations.
- **Longitudinal Health Tracking**: Integrate with health monitoring systems to adapt recommendations based on changing health metrics.
- **Personalized Nutritional Goals**: Develop algorithms that optimize recommendations to help users meet specific nutritional targets over time.

9.3.2 Enhanced Model Architecture

- **Deep Learning Integration**: Explore transformer-based models for improved understanding of meal descriptions and user preferences.
- **Multi-Modal Recommendations**: Incorporate image-based meal recognition to enhance the recommendation process.
- **Reinforcement Learning**: Implement reinforcement learning techniques to optimize long-term user health outcomes rather than just immediate engagement.

9.3.3 System Expansion

- **Meal Planning Integration**: Extend the system to generate complete meal plans that balance nutrition across multiple days.
- **Recipe Modification**: Develop capabilities to suggest personalized modifications to recipes based on dietary restrictions.
- **Social Recommendation Features**: Implement group recommendation capabilities for households with diverse dietary needs.
- **Explainable AI Components**: Enhance the system with clear explanations for why specific meals are recommended, increasing user trust and education.

These future directions aim to address current limitations while expanding the system's capabilities to provide more comprehensive, personalized nutrition guidance that adapts to users' evolving needs and preferences.

9.4 Ethical Considerations

As we continue to develop and refine NutriBuddy, several ethical considerations must be addressed:

- **Data Privacy**: Ensure robust protection of sensitive health information while maintaining recommendation quality.

- **Recommendation Diversity**: Prevent algorithmic bias that might limit exposure to certain food categories or cultures.
- **Health Responsibility**: Clearly communicate that recommendations complement but do not replace professional medical advice.

These ethical considerations will guide future development to ensure NutriBuddy remains a responsible and beneficial tool for improving nutritional health.

10 Conclusion

10.1 Key Takeaways

The NutriBuddy recommendation system successfully demonstrates the effectiveness of a hybrid approach to personalized nutrition recommendations. Our experimentation yielded several important findings:

- **Hybrid Model Superiority**: The combination of content-based and collaborative filtering (60%/40% weighting) consistently outperformed individual approaches across all evaluation metrics, with improvements of 5-7% in precision and recall.
- **Health-Conscious Personalization**: The system effectively balances personalization with health considerations, successfully matching meals to specific health conditions while maintaining recommendation diversity.
- **Dynamic Re-Ranking Impact**: Real-time adaptation to user feedback significantly improved engagement metrics, with a 48% increase in click-through rates and a 61% increase in purchase rates.
- **Cold-Start Problem Mitigation**: Enhanced content-based filtering and LLM-processed health data reduced the recommendation quality gap for new users by 32%.

10.2 Contribution to Project Objectives

These results directly address our primary project objectives:

- **Personalized Nutrition**: The system successfully provides tailored meal recommendations based on individual preferences, dietary restrictions, and health conditions.
- **Health Alignment**: Recommendations are effectively matched with specific health conditions, supporting users in making beneficial dietary choices.
- **Adaptive Learning**: The dynamic re-ranking mechanism ensures the system continuously improves based on user feedback and changing preferences.
- **Accessibility**: The intuitive dashboard interface makes personalized nutrition guidance accessible to users regardless of nutritional expertise.

10.3 Future Refinements

Based on our findings, we have identified several promising directions for further refinement:

- **Temporal Context Integration**: Incorporating time-of-day and seasonal factors to further personalize recommendations.
- **Meal Plan Generation**: Extending the system to create complete, nutritionally balanced meal plans across multiple days.
- **Explainable Recommendations**: Enhancing transparency by providing clear explanations for why specific meals are recommended.
- **Multi-Modal Input**: Exploring image-based meal recognition to enhance the recommendation process and user experience.

The NutriBuddy system represents a significant step forward in personalized nutrition technology, effectively combining data-driven recommendations with health-conscious guidance to support better dietary choices.

11 Reproducibility

11.1 Code Repository

The complete codebase for the NutriBuddy system is available on GitHub:
https://github.com/VenkatSR-14/nutribuddy

The repository is organized as follows:

- `/data`: Contains data preprocessing scripts and sample datasets
- `/models`: Implementation of content-based, collaborative filtering, and hybrid recommendation models
- `/experiments`: Evaluation scripts and experiment configurations
- `/app`: Dashboard implementation and user interface components
- `/docs`: Documentation and setup instructions

11.2 Dataset Access

The primary dataset used in this project is the Fitness Recommender Dataset, available on Kaggle: https://www.kaggle.com/datasets/venkyy123/fitness-recommender-dataset/data

This dataset was augmented with:

- Synthetic user profiles and interaction data generated using LLMs
- Enhanced nutritional information and health condition mappings
- User interaction logs collected during system development

11.3 Dependencies and Environment

The NutriBuddy system requires the following dependencies:

- **Python**: 3.9.x
- **Data Processing**: pandas (1.5.3), numpy (1.24.3)
- **Machine Learning**: scikit-learn (1.2.2), tensorflow (2.12.0)
- **NLP Processing**: nltk (3.8.1), spacy (3.5.3)
- **Web Framework**: FastAPI (2.3.2), React (18.2.0)
- **Database**: PostgreSQL
- **Visualization**: matplotlib (3.7.1), seaborn (0.12.2)

A complete environment can be recreated using the `requirements.txt` file in the repository. Docker configuration files are also provided for containerized deployment.

11.4 Experiment Reproduction

To reproduce our experiments:

- Clone the repository: `git clone https://github.com/VenkatSR-14/nutribuddy.git`
- Ensure Docker and Docker Compose are installed on your system
- Navigate to the project directory: `cd nutribuddy`
- Build and start the containerized environment: `docker-compose up --build`
- Access the dashboard through your browser at `http://localhost:5000`

The containerized setup automatically handles all dependencies, database initialization, and application deployment. This ensures a consistent environment across different systems and eliminates potential configuration issues.For development purposes, detailed setup instructions for local installation are available in the repository's `README.md` file.

12 References

1. M. Ilham Hafizha and Z. K. A. Baizal, "Diet and Physical Exercise Recommendation System Using a Combination of K-Means and Random Forest," Ind. Journal on Computing, vol. 9, no. 2, pp. 83-96, Aug. 2024. doi:10.34818/indojc.2024.9.2.959.

2. B. M. Hassan and S. M. Elagamy, "Personalized medical recommendation system with machine learning," Neural Computing and Applications, Springer Nature, 2025. doi:10.1007/s00521-024-10916-6.

3. S. Nagati, J. Chaitanya, B. S. Rani, and R. Rohan, ”Diet Recommendation Systems Using Machine Learning Approaches, Challenges, and Future Directions,” International Research Journal of Education and Technology, vol. 6, no. 12, pp. 1872-1876, Dec. 2024.

4. B. Boudaa, A. Hammadi, and K. Akermi, ”A Diet Recommendation System Using TF-IDF and Extra Trees Algorithm,” International Journal of Electrical and Electronics Engineering and Informatics, vol. 12, no. 4, 2024. doi:10.52549/ijeei.v12i4.5809.

5. G. Qiao, D. Zhang, N. Zhang, X. Shen, X. Jiao, W. Lu, D. Fan, J. Zhao, H. Zhang, W. Chen, and J. Zhu, ”Food Recommendation Towards Personalized Wellbeing,” Trends in Food Science & Technology, vol. 156, Feb. 2025, 104877. doi:10.1016/j.tifs.2025.104877.